
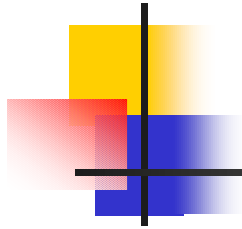


# Advanced Sockets API for IPv6



yhmiu

RFC 2292  
RFC 1883



# Outline

---

- IPv6 Raw sockets
- Ancillary Data
- Options described using ancillary data
  - Packet information
  - Hop-By-Hop Options
  - Destination Options
  - Routing Header Option



# IPv6 Raw sockets

---

- Raw sockets bypass the transport layer ( TCP or UDP )
- All fields in the IPv6 header that an application might want to change can be modified using ancillary data and/or socket options by the application for **output**
- All fields in a received IPv6 header and all extension headers are also made available to the application as ancillary data on **input**



# Example of ICMPv6 Raw socket

---

- Calculate and insert the ICMPv6 checksum for ICMPv6 raw socket

```
fd=socket(PF_INET6,SOCK_RAW,IPPROTO_ICMPV6);  
int offset=2;  
setsockopt(fd,IPPROTO_IPV6,IPV6_CHECKSUM,&offset,sizeof(offset));
```



# Ancillary Data (1)

---

- Be used to exchange the following optional information between the AP and kernel
  1. The send/receive interface and source/destination address
  2. The hop limit
  3. Next hop address
  4. Hop-By-Hop options
  5. Destination options
  6. Routing header



## Ancillary Data (2)

---

- 4.2BSD allowed file descriptors to be transferred between separate processes across a UNIX domain socket using `sendmsg()` and `recvmsg()` functions
- Two member of `msg_hdr` structure will be used
  - `msg_control`
  - `msg_controllen`



## Ancillary Data (3)

---

- Ancillary data object
  - Defined by the `cmsghdr` structure
- Functions that operate on the ancillary data objects
  - `GMSG_FIRSTHDR()`
  - `GMSG_NXTHDR()`
  - `GMSG_DATA()`
  - `GMSG_SPACE()`
  - `GMSG_LEN()`



# The msghdr structure

---

```
struct msghdr {  
    void      *msg_name;      /* ptr to socket address structure */  
    socklen_t  msg_namelen;    /* size of socket address structure */  
    struct iovec *msg_iov;     /* scatter/gather array */  
    size_t     msg_iovlen;     /* # elements in msg_iov */  
    void      *msg_control;    /* ancillary data */  
    socklen_t  msg_controllen; /* ancillary data buffer length */  
    int        msg_flags;      /* flags on received message */  
};
```



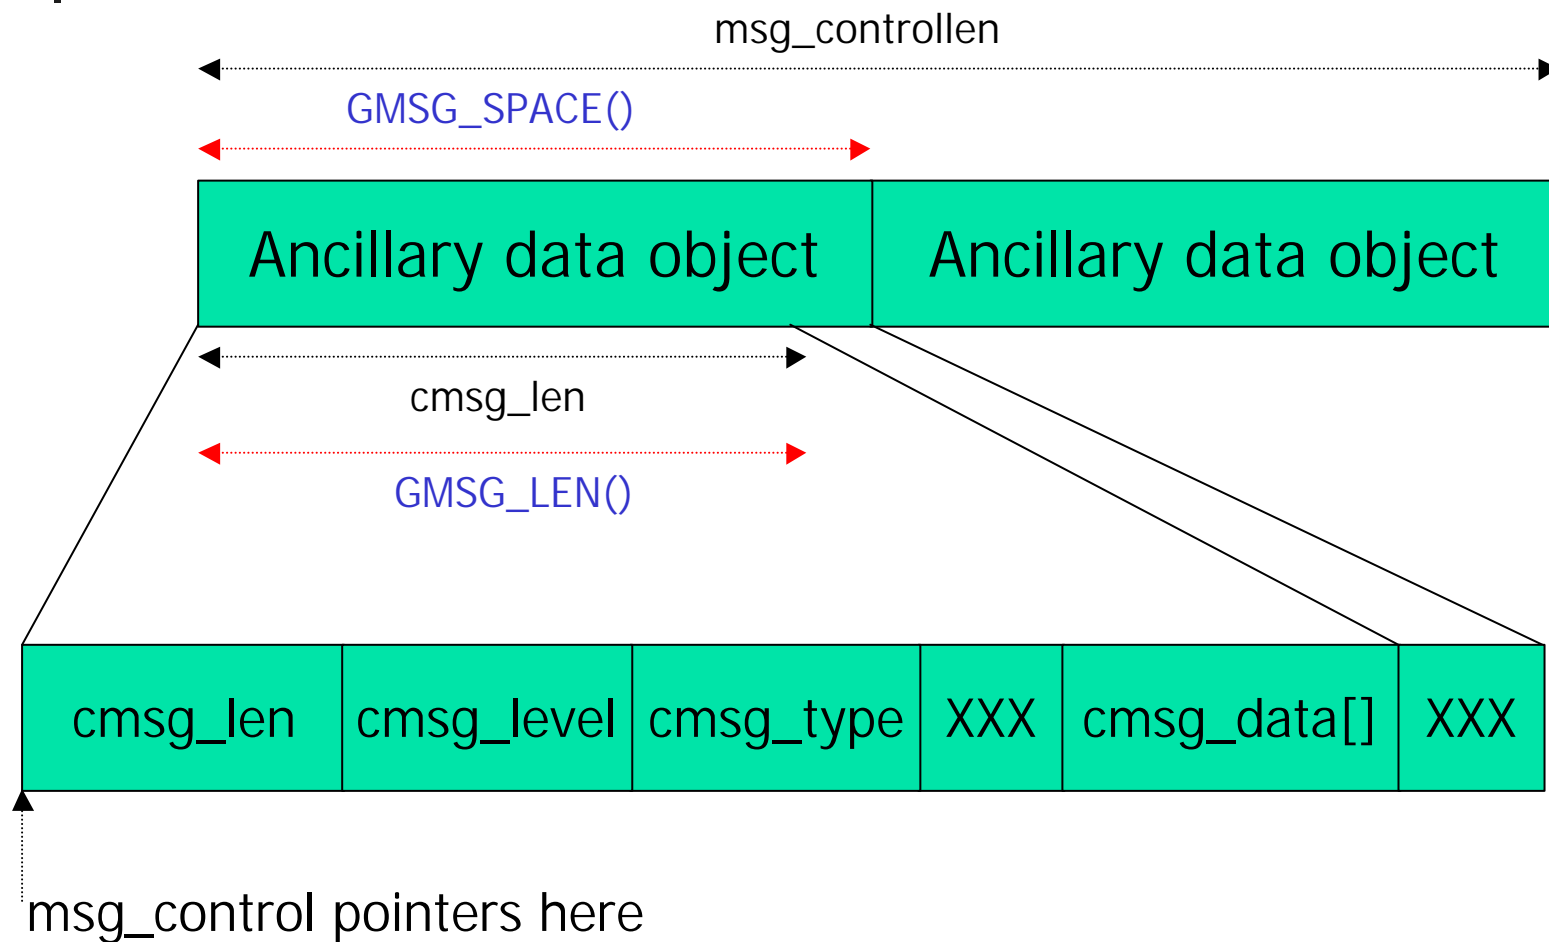


# The cmsghdr structure

---

```
struct cmsghdr {  
    socklen_t  cmsg_len; /* #bytes, including this header */  
    int        cmsg_level; /* originating protocol */  
    int        cmsg_type; /* protocol-specific type */  
    /* followed by unsigned char cmsg_data[]; */  
};
```

# Ancillary Data (4)





# Ancillary data object macros(1)

---

- **GMSG\_FIRSTHDR()**
  - `struct cmsghdr *CMSG_FIRSTHDR(const struct msghdr *mhdr);`
  - Return a pointer to the first cmsghdr structure in the msghdr structure pointed to by mhdr
- **GMSG\_NXTHDR()**
  - `struct cmsghdr *CMSG_NXTHDR(const struct msghdr *mhdr, const struct cmsghdr *cmsg);`
  - Return a pointer to the cmsghdr structure describing the next ancillary data object
  - mhdr is a pointer to a msghdr structure and cmsg is a pointer to a cmsghdr structure
  - If there is not another ancillary data object, the return value is NULL



## Ancillary data object macros(2)

---

- `GMSG_DATA()`
  - `unsigned char *CMSG_DATA(const struct cmsghdr *cmsg);`
  - Return a pointer to the data (`cmsg_data[]`) following a `cmsghdr` structure
- `GMSG_SPACE()`
  - `unsigned int CMSG_SPACE(unsigned int length);`
  - returns the space required by the object and its `cmsghdr` structure, including any padding needed to satisfy alignment requirements



## Ancillary data object macros(3)

---

- CMSG\_LEN()
  - unsigned int CMSG\_LEN(unsigned int length);
  - returns the value to store in the cmsg\_len member of the cmsghdr structure, taking into account any padding needed to satisfy alignment requirements



# Options described using ancillary data

---

1. The application must call `setsockopt()` to turn on the corresponding flag, then to receive any of this optional information
2. When any of these options are enabled, the corresponding data is returned as control information by `recvmsg()`, as one or more ancillary data objects
3. 在傳送端，the application just calls `sendmsg()` and specifies one or more ancillary data objects as control information.並不需要像接收端一樣要call `setsockopt()`



# cmsghdr fields of the ancillary data objects

cmsg_level	cmsg_type	cmsg_data[]
IPPROTO_IP6	IPV6_PKTINFO	in6_pktinfo structure
IPPROTO_IP6	IPV6_HOPLIMIT	int
IPPROTO_IP6	IPV6_NEXTHOP	Socket address structure
IPPROTO_IP6	IPV6_HOPOPTS	implementation dependent
IPPROTO_IP6	IPV6_DSTOPTS	implementation dependent
IPPROTO_IP6	IPV6_RTHDR	implementation dependent



# Routing header (1)

---

- IPv6 current define only Type 0 Routing header
  - This type supports up to 23 intermediate nodes
  - Each hop is defined as a strict or loose hop





## Routing header (2)

---

- 4 functions build a Routing header
  - `inet6_rthdr_space()`
    - Return #bytes required for ancillary data
  - `inet6_rthdr_init()`
    - Initialize ancillary data for Routing header
  - `inet6_rthdr_add()`
    - Add IPv6 address & flags to Routing header
  - `inet6_rthdr_lasthop()`
    - Specify the flags for the final hop



## Routing header (3)

---

- 4 functions deal with a returned Routing header
  - `inet6_rthdr_reverse()`
    - Reverse a Routing header
  - `inet6_rthdr_segments()`
    - Return #segments in a Routing header
  - `inet6_rthdr_getaddr()`
    - Fetch one address from a Routing header
  - `inet6_rthdr_getflags()`
    - Fetch one flag from a Routing header
- 以上8個function皆被定義於<netinet6/in6.h>



## Routing header (4)

---

- In order to receive a Routing header, the application must enable the IPV6\_RTHDR option

```
int on = 1;
```

```
setsockopt(fd, IPPROTO_IPV6, IPV6_RTHDR, &on, sizeof(on));
```

- To send a Routing header the application just specifies it as ancillary data in a call to `sendmsg()`




# 補充

---

- The following constants are defined in the `<netinet6/in6.h>` header


```
#define IPV6_RTHDR_LOOSE    0 /* this hop need not be a neighbor */  
#define IPV6_RTHDR_STRICT  1 /* this hop must be a neighbor */  
#define IPV6_RTHDR_TYPE_0  0 /* IPv6 Routing header type 0 */
```

# Routing header example



src:	*	S	S	S	S	S
dst:	D	I1	I2	I3	D	D
A[1]:	I1	I2	I1	I1	I1	I1
A[2]:	I2	I3	I3	I2	I2	I2
A[3]:	I3	D	D	D	I3	I3
#seg:	3	3	2	1	0	3

check:    f0                      f1                      f2                      f3



## S端AP在call sendmsg()前該做的事


```
void *ptr;
struct msghdr msg;
struct cmsghdr *cmsgptr;
struct sockaddr_in6 I1, I2, I3, D;
unsigned int f0, f1, f2, f3;

ptr = malloc(inet6_rthdr_space(IPV6_RTHDR_TYPE_0, 3));
cmsgptr = inet6_rthdr_init(ptr, IPV6_RTHDR_TYPE_0);

inet6_rthdr_add(cmsgptr, &I1.sin6_addr, f0);
inet6_rthdr_add(cmsgptr, &I2.sin6_addr, f1);
inet6_rthdr_add(cmsgptr, &I3.sin6_addr, f2);
inet6_rthdr_lasthop(cmsgptr, f3);

msg.msg_control = ptr;
msg.msg_controllen = cmsgptr->cmsg_len;

/* finish filling in msg{}, msg_name = D */
/* call sendmsg() */
```




call `inet6_rthdr_init()` initializes the ancillary data object to contain a Type 0 Routing header

0

31

cmsg_len=20			
cmsg_level=IPPRTO_IPV6			
cmsg_type=IPV6_RTHDR			
Next Header	Hdr Ext Len=0	Routing Type=0	Seg Left=0
Reserved	Strict/Loose Bit Map		




The first call to `inet6_rthdr_add()` adds I1 to the list

0

31

cmsg_len=36			
cmsg_level=IPPRTO_IPV6			
cmsg_type=IPV6_RTHDR			
Next Header	Hdr Ext Len=2	Routing Type=0	Seg Left=1
Reserved	X	Strict/Loose Bit Map	
Address[1]=I1			






The next call to `inet6_rthdr_add()` adds I2 to the list

0

31

cmsg_len=52			
cmsg_level=IPPRTO_IPV6			
cmsg_type=IPV6_RTHDR			
Next Header	Hdr Ext Len=4		Routing Type=0
Seg Left=2			
Reserved	X	X	Strict/Loose Bit Map
Address[1]=I1			
Address[2]=I2			




The last call to `inet6_rthdr_add()` adds I3 to the list

0

31

cmsg_len=68				
cmsg_level=IPPRTO_IPV6				
cmsg_type=IPV6_RTHDR				
Next Header	Hdr Ext Len=6		Routing Type=0	Seg Left=3
Reserved	X	X	X	Strict/Loose Bit Map
Address[1]=I1				
Address[2]=I2				
Address[3]=I3				



the call to `inet6_rthdr_lasthop()` sets the next bit of the  
Strict/Loose Bit Map to the value specified by `f3`

0

31

cmsg_len=68											
cmsg_level=IPPRTO_IPV6											
cmsg_type=IPV6_RTHDR											
Next Header		Hdr Ext Len=6				Routing Type=0		Seg Left=3			
Reserved		X	X	X	X	Strict/Loose Bit Map					
Address[1]=I1											
Address[2]=I2											
Address[3]=I3											